

# Next-Gen Model Risk Management

## High-Fidelity Document Intelligence

### Transforming PDF Analysis with AI-Powered Extraction

Alexander Tsoskounoglou

Associate, Model Risk Management — Interest Rate Derivatives (MRM IRD)  
Morgan Stanley, Budapest

February 2026

## Alexander Tsoskounoglou

- **BSc Computer Science** — solid foundation in software engineering & algorithms
- **MSc Computational Finance** (University of Amsterdam) — quantitative modelling, stochastic calculus & numerical methods
- **Quantitative Analyst** (1 year, hedge fund) — developed & deployed neural-network-based trading strategies in production
- **Quantitative Researcher** (1 year, University of Amsterdam) — led two research projects on Deep Reinforcement Learning for hedging & pricing exotic derivatives

# The Current Friction: A Wall of PDFs

## The MRM Document Landscape

- **Thousands** of model documents in the Model Context System
- Bermudian Swaptions, SABR Calibrations, CVA/XVA frameworks...
- Dense with:
  - Multi-line integral equations & SDEs
  - Stochastic calculus (Itô, martingales)
  - Nested tables of Greeks & sensitivities
  - **Figures, charts & diagrams** — calibration surfaces, payoff diagrams

## Why “Copy-Paste into an LLM” Fails

- ① **PDF  $\neq$  Text** — PDFs store glyphs, not semantics
- ② **Math is destroyed** — integrals → dashes, Greek → ASCII
- ③ **Tables collapse** — columns merge into gibberish
- ④ **Figures vanish** — images are silently dropped
- ⑤ **LLM hallucination** — garbage in ⇒ *confident* garbage out

⚠ This is not a minor inconvenience.  
It is a **model risk**.

# The Technical Gap: “If the LLM Can’t Read the Math...”

Source: *Markov Functional Market Model* (Oxford, 2008) — HJM Drift Condition & Bond Pricing

## ✗ Standard Text Extraction

```
P(t, T) = EQ exp rs ds ???| t  
      - t F  
      - T  
alpha(t, T) = sigma(t, T) · t  
t sigma(t, s) ds ???? 0 <= t <= T.  
P(t, T) = exp A(t, T) B(t, T) rt  
At betaB + [1][= 0][,]  
      - 2 [delta] [() [t] ()] [B] [2]
```

- ✗ Integrals → dashes, measures → gone
- ✗ Greek letters → English words
- ✗ Square brackets are OCR artifacts
- ✗ Completely unusable for validation

## ✓ High-Fidelity Pipeline (Marker/Surya)

$$\begin{aligned} P(t, T) &= \mathbb{E}_{\mathbb{Q}} \left[ \exp \left( - \int_t^T r_s \, ds \right) \middle| \mathcal{F}_t \right] \\ \alpha(t, T) &= \sigma(t, T) \cdot \int_t^T \sigma(t, s) \, ds \quad \forall 0 \leq t \leq T \\ P(t, T) &= \exp(A(t, T) - B(t, T) r_t) \end{aligned}$$

- ✓ Integral notation & bounds preserved
- ✓ Measure  $\mathbb{Q}$ , filtration  $\mathcal{F}_t$  correct
- ✓ Greek letters  $\alpha, \beta, \gamma, \delta$  intact
- ✓ Machine-readable & LLM-ready

# Real-World Failure: Longstaff-Schwartz Continuation Value

**Source:** Same document — Longstaff-Schwartz algorithm (Eq. 4.1).

## ✗ Standard Extraction

```
C(w; tk) = EQ[????? j=k+1 N  
exp(- ?? rj) I(w, tj; tk, T)] ???Ftk
```

LLM Interpretation:

- ✗ Summation sign → “?????”
- ✗ Integral replaced by “– ??”
- ✗ Conditional  $\mathcal{F}_{t_k}$  → “???Ftk”

An LLM will hallucinate the missing structure.

## ✓ High-Fidelity Output

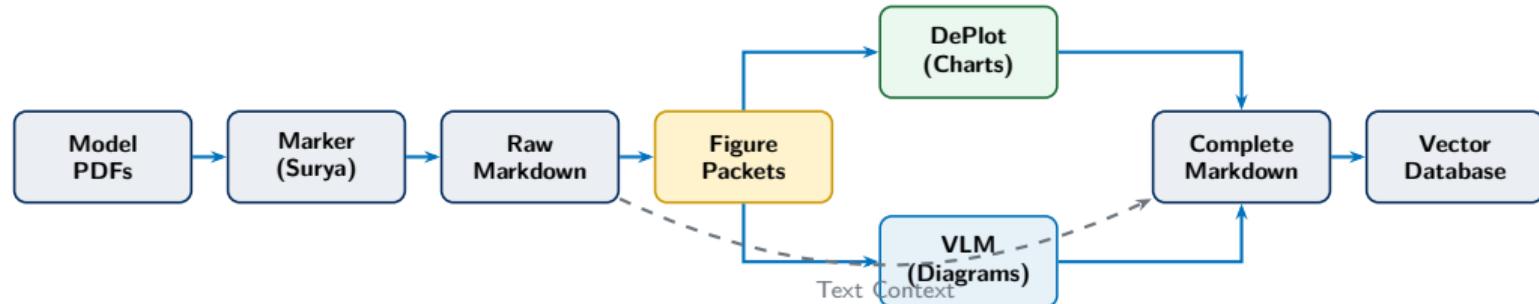
Continuation value at exercise time  $t_k$ :

$$C(\omega; t_k) = \mathbb{E}_{\mathbb{Q}} \left[ \sum_{j=k+1}^K \exp \left( - \int_{t_k}^{t_j} r(s) ds \right) I(\omega, t_j; t_k, T) \right]$$

- ✓ Summation, integral, conditional correct
- ✓ Validator can cross-reference with code

The equation we use for Bermudan swaption validation.

# Proposed Solution: End-to-End Document Intelligence



## How It Works

- **Figure Packets:** Context-aware extraction (Image + Caption + In-text mentions).
- **Chart Data:** DePlot converts plots to structured tables (not just pixels).
- **Diagrams:** Multimodal LLM provides structured JSON interpretation.

## The End-State Vision

- **Conversational Querying:** Expert answers with equation references.
- **High-Fidelity Context:** LLM sees the data behind the charts.
- **Accelerated Workflows:** Clean context for testing & learning.

# Research Summary: Key Technologies

Tool	Role	Strength	Hardware	Key Detail
<b>Marker</b> (+ Surya)	<b>PDF → Mark-down</b> (core pipeline)	✓ Best-in-class: 95.7 score	GPU or CPU	25 pages/sec on GPU; Surya is the OCR engine.
<b>DePlot</b> (Google)	<b>Chart → Table</b> (data extraction)	✓ Linearizes plots to text	GPU	Converts chart images into data tables for reasoning.
<b>Nemotron</b> (NVIDIA)	1. <b>Figure/Image AI</b> 2. <b>Quant Risk Chat</b>	✓ 1M context window; SOTA open-model	Cost-effective GPUs (A10)	Process entire PVRMs in one pass. Ideal for both visual reasoning & risk analysis conversations.

## Strategy

**Marker** for text/math. **DePlot** for charts. **Nemotron** for deep reasoning and diagrams.

# Why This Matters: Business Value

## Speed of Validation

- **Current:** reading 50–400 page PVRM, manually locating equations — *days/weeks*
- **Proposed:** natural language query in seconds.

Task	Current	Proposed
Locate equation	15–30 min	<1 min
Cross-ref docs	2–4 hrs	5–10 min
Review documentation	1–2 weeks	~3–5 days
Write new tests	Weeks	Days

## Broader Impact

- **Reduce model risk:** automated consistency checks between spec and code
- **Replicate & extend tests:** brainstorm new validation tests from the math
- **Assist in Extending & Merging PVRMs:** Help analysts identify shared assumptions/gaps
- **Onboarding:** new analysts query models immediately
- **Audit trail:** logged queries for compliance

# Implementation Phases

## Phase 1: Proof of Concept

- ✓ Convert IR\_Swaption\_Bermudan PVRM (12 PDFs)
- ✓ Validate math fidelity
- ✓ LLM Context (Copilot)
- ✓ Collect **team feedback**
- ✓ Set accuracy & performance standards

≈ 4–8 weeks

## Phase 2: Pilot Vector DB

- ▶ Build vector DB for pilot team
- ▶ Nemotron for Figure AI
- ▶ Iterate on infrastructure
- ▶ Measure productivity gains
- ▶ Infra support needed

≈ 3–5 months

## Phase 3: MRM Rollout

- ▶ Deploy across MRM
- ▶ Multi-PVRM cross-ref
- ▶ Analyst Web Interface
- ▶ Training & Docs
- ▶ Model updates

≈ 6–12 months

Month 1–2

Month 3–7

Month 8–14+

# Technical Architecture

## 1. Ingestion Pipeline

- **PDF Parsing:** Marker (**Surya**) for layout/math
- **Figure Packets:** Image + Caption + Context
- **Routing:** DePlot (Charts) & VLM (Diagrams)
- **Quality Gates:** Consistency checks (e.g. caption match)
- **Result:** Full Markdown with structured data

## 2. Dual-Store Storage

- **Vector DB:** Equation-aware chunks
- **Doc Store:** Faithful Markdown context

## 3. Query & Retrieval Flow

- **Retriever:** Fetches ranked chunks + full context
- **LLM Engine:** Nemotron generates answers
- **Analyst Interface:** Secure user access

## Security & Compliance

- **100% On-Premise:** No external API calls
- **Open Source Stack:** Fully auditable code
- **Access Control:** Respects team permissions
- **Audit Trails:** All queries are logged

# Next Steps & Requirements

## Phase 1 Actions

- ① **Approval** to build local prototype
- ② Convert **PVRM IR\_Swaption\_Bermudan** (12 PDFs)
- ③ Validate fidelity against original documents
- ④ Define math fidelity score & criteria
- ⑤ Pilot AI assistants & collect feedback

## Resource Requirements

Resource	Specification
GPU	A10 or A100 (recommended)
CPU Fallback	Possible (slower)
Storage	~50 GB
Time	4–8 weeks (Phase 1)
Team	1 engineer

## The Ask

Approval to allocate compute access and begin Phase 1 prototype on the Bermudan swaption PVRM.

# Thank You

## Questions?

`alexander.tsoskounoglou@morganstanley.com`

**MRM IRD — Budapest**

*"The quality of AI output is bounded by the quality of its input."*

# Appendix: Under the Hood

**Equation:** Longstaff-Schwartz Continuation Value (Eq. 4.1)

## 1. Original PDF

$$C(\omega; t_k) = \mathbb{E}_{\mathbb{Q}} \left[ \sum_{j=k+1}^K \exp \left( - \int_{t_k}^{t_j} r(\omega, s) ds \right) I(\omega, t_j; t_k, T) \middle| \mathcal{F}_{t_k} \right]$$

## 2. Standard Text Extraction (pymupdf4llm)

```
C(w; tk) = EQ[????? j=k+1 N exp(- ?? rj) I(w, tj; tk, T)] ???Ftk
```

## 3. High-Fidelity Pipeline (Marker/Surya)

```
$$C(w; t_k) = \mathbb{E}_{\mathbb{Q}} \left[ \sum_{j=k+1}^K \exp \left( - \int_{t_k}^{t_j} r(w, s) ds \right) I(w, t_j; t_k, T) \dots \right]
```

# Appendix: Marker Benchmark Performance

Method	Time (sec)	Heuristic Score	LLM Score
Marker	<b>2.84</b>	<b>95.67</b>	<b>4.24</b>
Doclinc	3.70	86.71	3.70
Mathpix	6.36	86.43	4.16
LlamaParse	23.35	84.24	3.98

Benchmark data from Marker GitHub repository.

## Throughput

- Single-page serial: 2.84 s
- Batch mode on H100: **25 pages/sec**
- Can run on **A10, V100, MPS, or CPU**

## By Document Type (Heuristic)

- |                    |       |
|--------------------|-------|
| Scientific paper   | 96.67 |
| Financial document | 95.37 |
| Legal document     | 96.69 |

**Financial documents** score 95.4/100.

# Appendix: Cost & Productivity Estimation

## Cost per Query

**Platform:** Microsoft Azure (MS strategic partner)

Model	In/1M	Out/1M
Nemotron-3-Nano (open-source)	\$0.05	\$0.20
Gemini 3 Pro (closed, SOTA)	\$2.00	\$12.00

**Typical query:** ~4k input + ~1.5k output tokens

- ▶ Nemotron-3-Nano: **\$0.0005**/query
- ▶ Gemini 3 Pro: **\$0.026**/query

## Monthly Cost Projection (per analyst)

**Est. queries:** 25–40/day × 22 days/month

Scenario	Q/day	Nano	Gemini 3
Conservative	25	\$0.28	\$14.30
Normal	35	\$0.39	\$20.02
Heavy use	50	\$0.55	\$28.60

**Team of 10 (Gemini 3): \$143–286/month**

## Estimated Efficiency Gains per PVRM Revalidation

Task	Current	Conservative	Normal	Optimistic
Read & digest PVRM docs	~16–24 hrs	–10% (14–22 h)	–20% (13–19 h)	–30% (11–17 h)
Locate equations/sections	~2–4 hrs	–15% (1.7–3.4 h)	–25% (1.5–3 h)	–40% (1.2–2.4 h)
Cross-reference documents	~4–8 hrs	–15% (3.4–6.8 h)	–30% (2.8–5.6 h)	–45% (2.2–4.4 h)
Interpret figures & tables	~2–4 hrs	–10% (1.8–3.6 h)	–20% (1.6–3.2 h)	–35% (1.3–2.6 h)
Draft validation tests	~40–80 hrs	–10% (36–72 h)	–20% (32–64 h)	–30% (28–56 h)
<b>Total per PVRM</b>	<b>~64–120 hrs</b>	<b>–12%</b>	<b>–22%</b>	<b>–35%</b>

Pricing: Azure AI Foundry (Feb 2026). Assumes RAG-based retrieval. Speed-ups assume high-fidelity extraction pipeline is deployed.